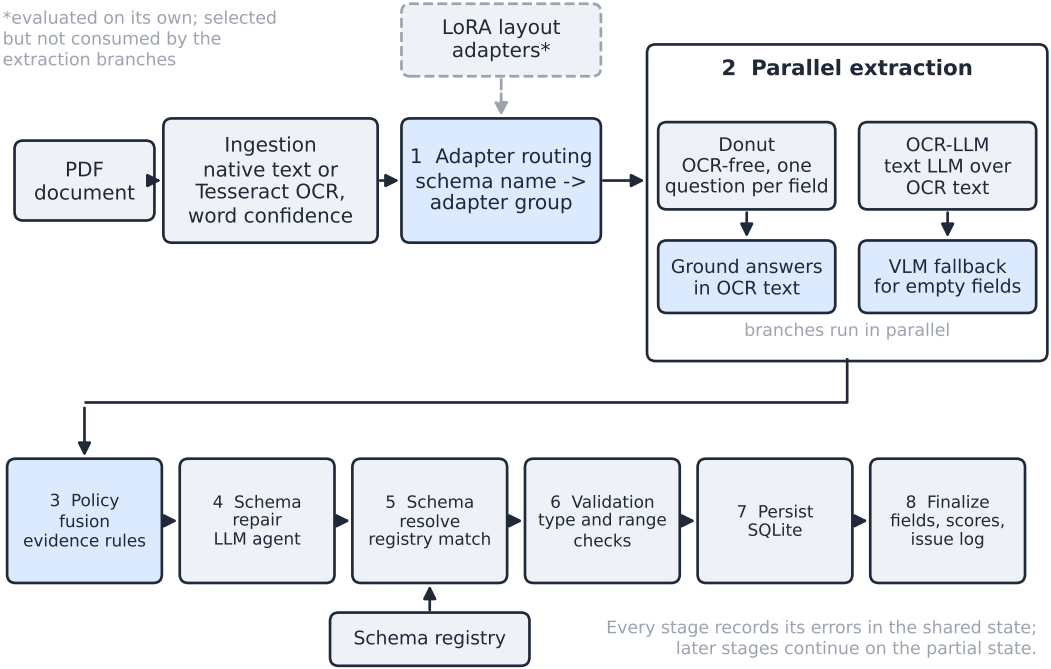
*evaluated on its own; selected but not consumed by the extraction branches
LoRA layout adapters*
PDF document
Ingestion native text or Tesseract OCR, word confidence
1 Adapter routing schema name -> adapter group
2 Parallel extraction
Donut OCR-free, one question per field
OCR-LLM text LLM over OCR text
Ground answers in OCR text
VLM fallback for empty fields
branches run in parallel
3 Policy fusion evidence rules
4 Schema repair LLM agent
5 Schema resolve registry match
6 Validation type and range checks
7 Persist SQLite
8 Finalize fields, scores, issue log
Schema registry
Every stage records its errors in the shared state; later stages continue on the partial state.